



Joining Kubernetes clusters with Submariner

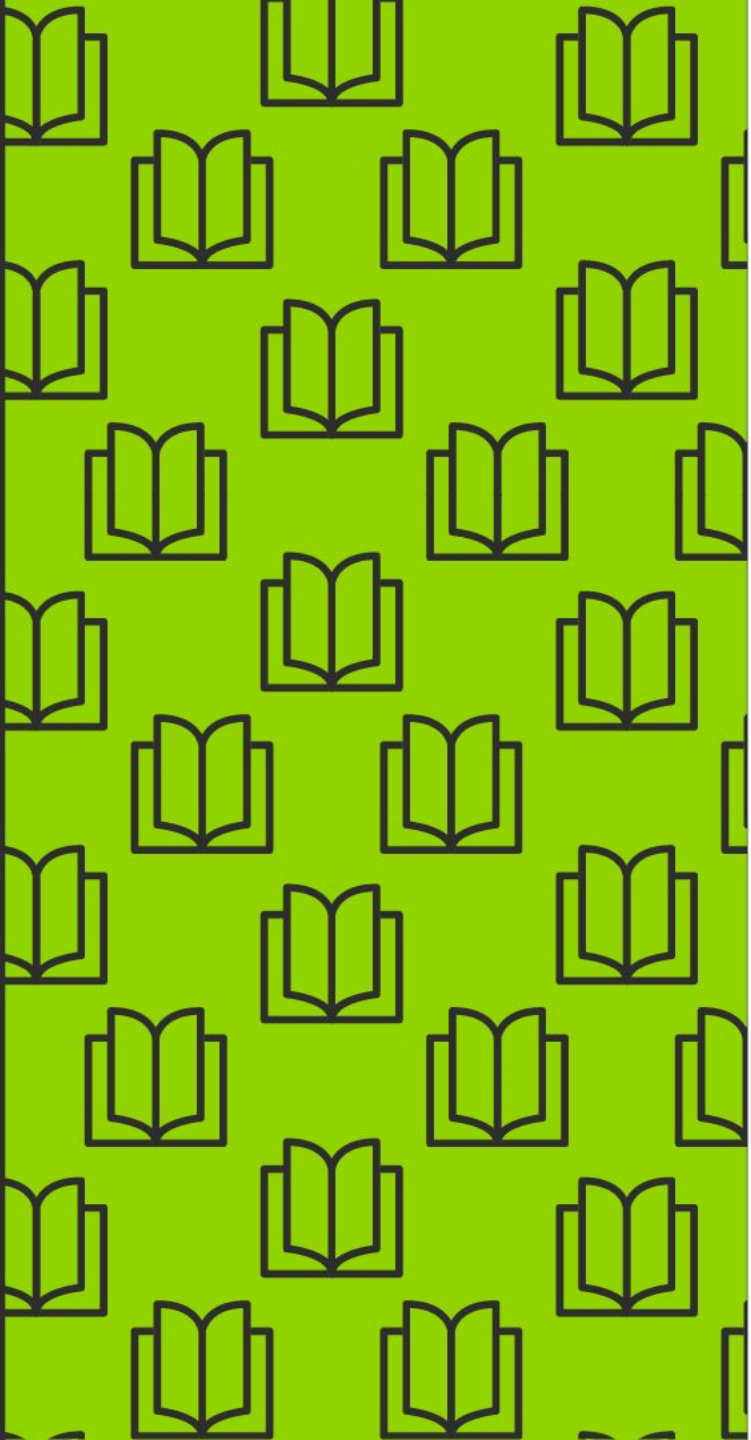
Community, research, product

Nir Yechiel
Engineering Manager
Red Hat

Daniel Bachar
MSc student
Reichman University (IDC Herzliya)

What we'll discuss today

- ▶ Why is multi-cluster connectivity important?
- ▶ Community: Submariner project
- ▶ Research: Optimized Service Discovery Across Clusters
- ▶ Product: Advanced Cluster Management for Kubernetes
- ▶ Resources




Why is
multiclust
networking
important?

Where is the growth in cluster deployments?



Small Scale Dev Teams

Managing clusters across
Dev/QA/Prod clusters




Medium Scale Organizations

Local retails with clusters
across 100s of locations



Large Scale Organizations

Global organizations with
100s of clusters, hosting
thousand of applications



Edge Scale Telco

100s of zones, 1000s of
clusters and nodes across
complex topologies

Reasons for deploying multiple clusters



Application
availability



Reduced
latency



Address industry
standards



Geopolitical data
residency guidelines



Disaster
recovery



Edge
deployments



CapEx
cost reduction



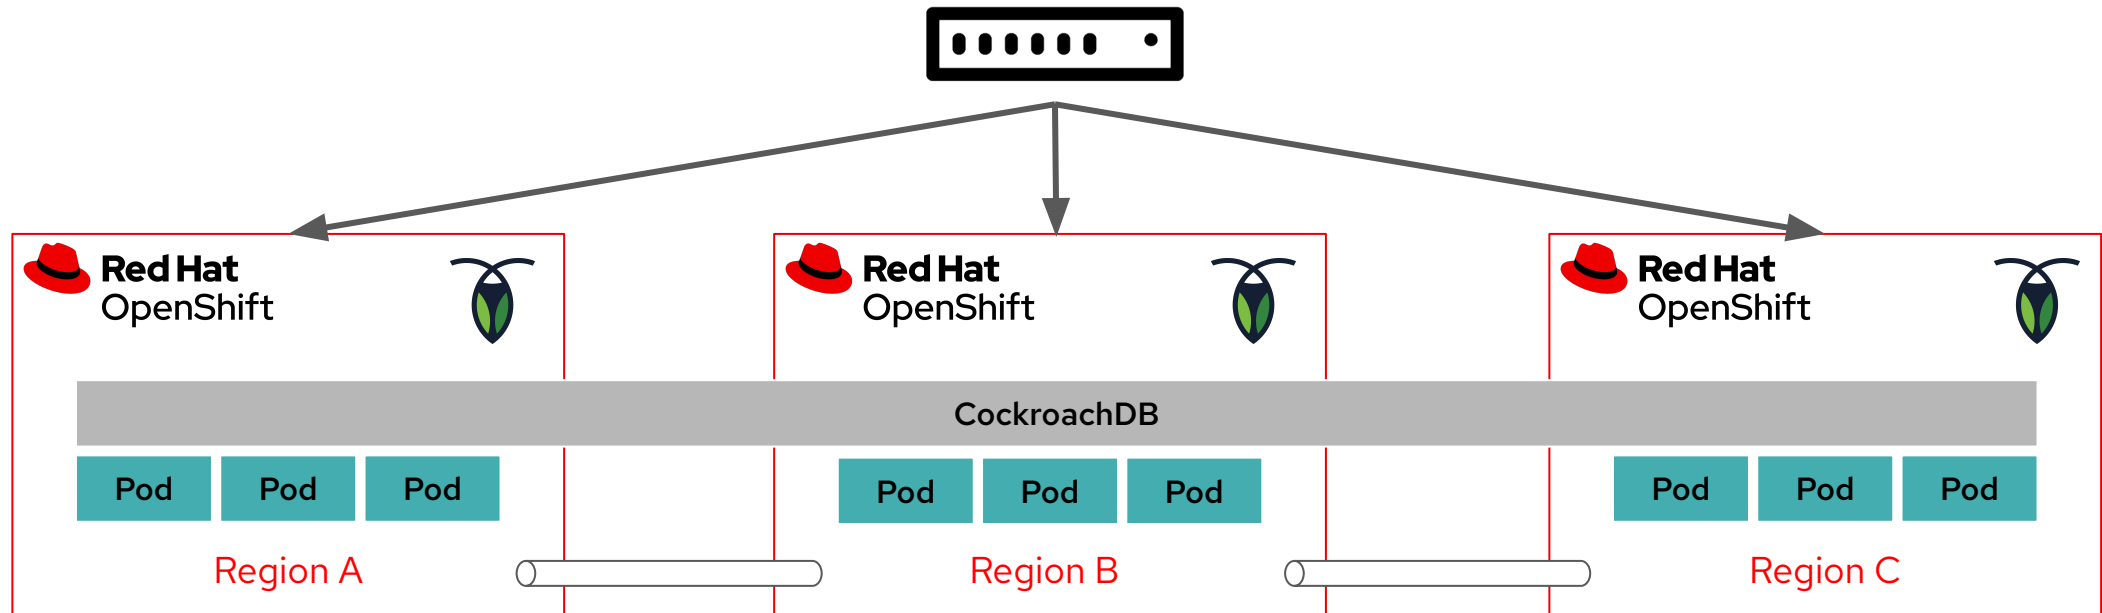
Avoid vendor
lock-in

Multicloud networking

- Multi-cluster *connectivity* is at the core of Red Hat's open hybrid cloud strategy and required for a wide range of use cases
- Our customers demand choice: a robust solution that works across different infrastructure providers and network (CNI) plug-ins
- Complement Red Hat's product portfolio:
 - Red Hat OpenShift
 - Advanced Cluster Management for Kubernetes (ACM)
 - Red Hat OpenShift Data Foundation (previously known as OCS)
 - Red Hat Service Mesh (Istio)

Example use case: Distributed Data


Multi-region CockroachDB cluster



- OpenShift clusters in multiple regions, with replicas of the same service running in each cluster
- Goal: keep data close to the user, while reducing latency and fault tolerance to improve user experience
- [Blog post](#) | [demo](#)

Example use case: Disaster Recovery (DR)

Cross-cluster volume replication with OpenShift Data Foundation



- Automated per application failover management through ACM
- Goal: protection against geographic disasters
- [Documentation](#) | [demo](#)



Community: Submariner project

Submariner project

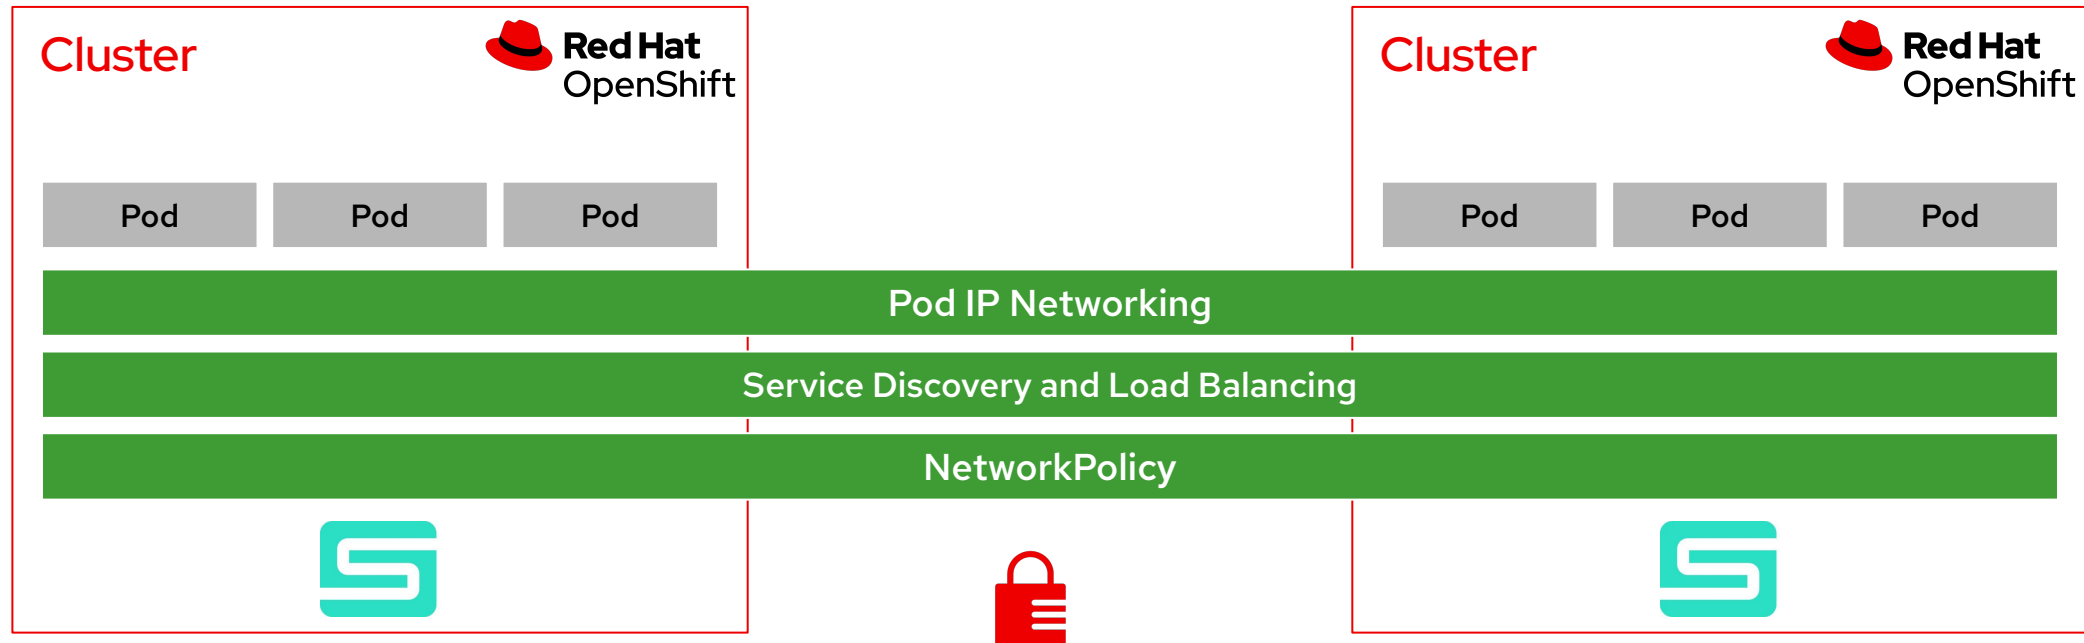
“Directly connect all your Kubernetes clusters, no matter where they are in the world”

<https://submariner.io>

<https://github.com/submariner-io>

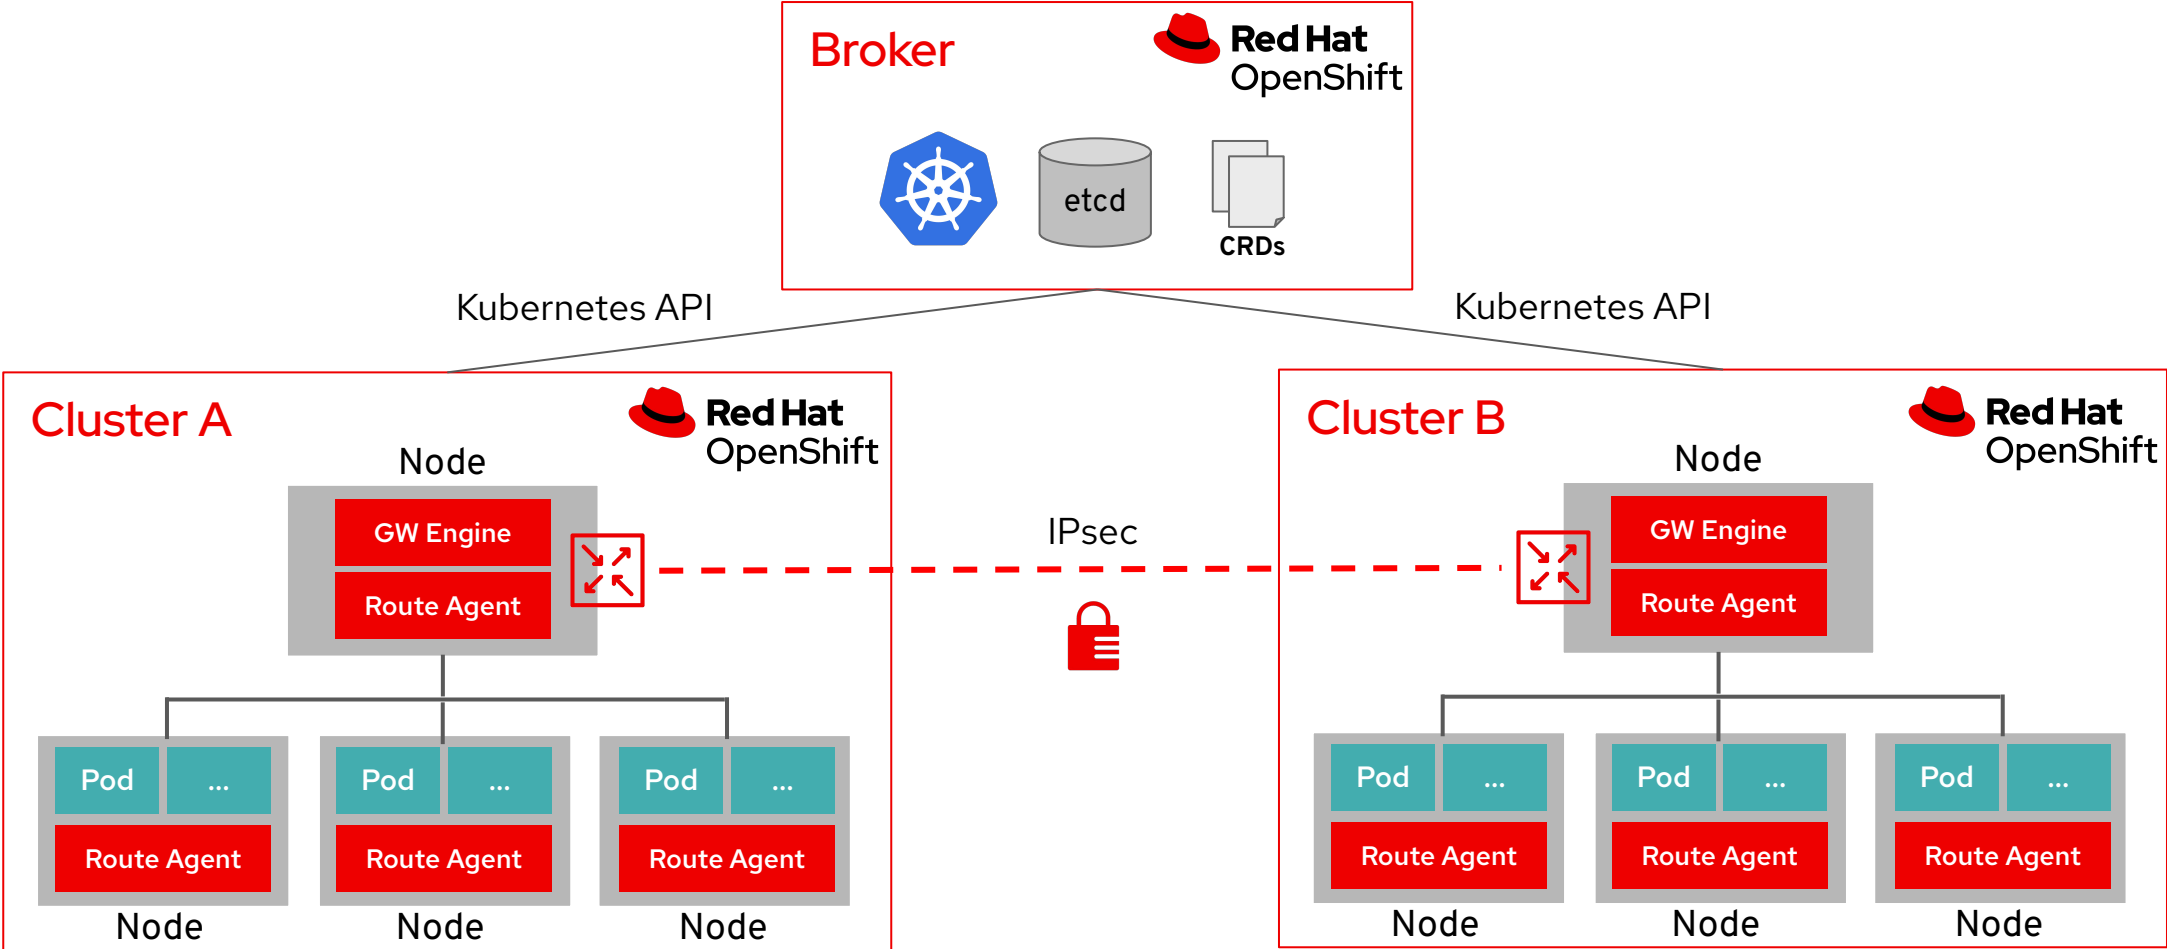


Kubernetes cluster networking with Submariner



- Different regions of the same public cloud provider
- Multiple public clouds
- Multiple on-prem sites
- Hybrid cloud, including a mix of on-prem and public cloud

Architecture



Benefits




Pod-to-pod and pod-to-service routing with native performance

Direct network tunnel to support any application on top; eliminate the need for proxies, external load-balancers or ingress gateways



Enhanced security

All traffic flow between clusters is encrypted using IPsec by default



Deploy services across clusters

Beyond connectivity, also address the challenge of cross-cluster service discovery and network policy (roadmap)



Extend existing OpenShift deployments

Compatible with different infrastructure providers and network (CNI) plugins; benefit the wider OpenShift ecosystem

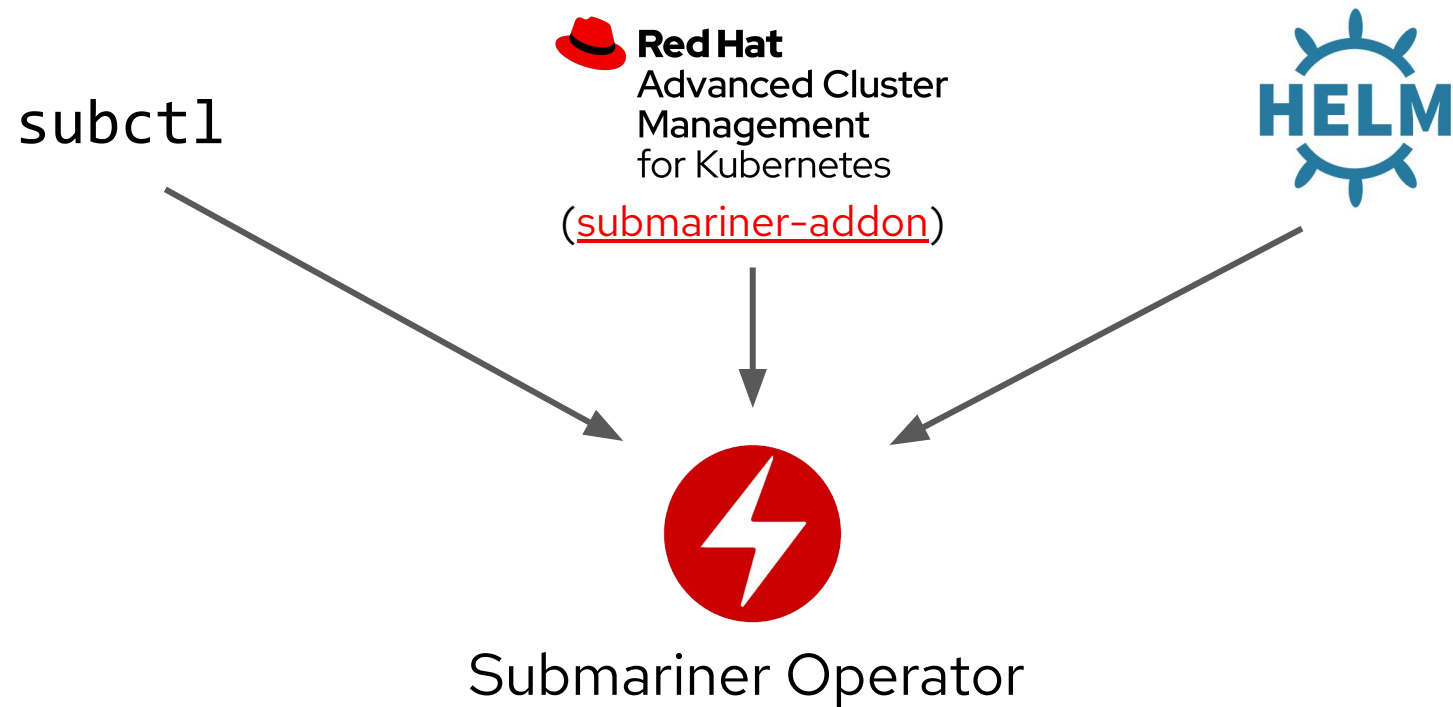
Usage

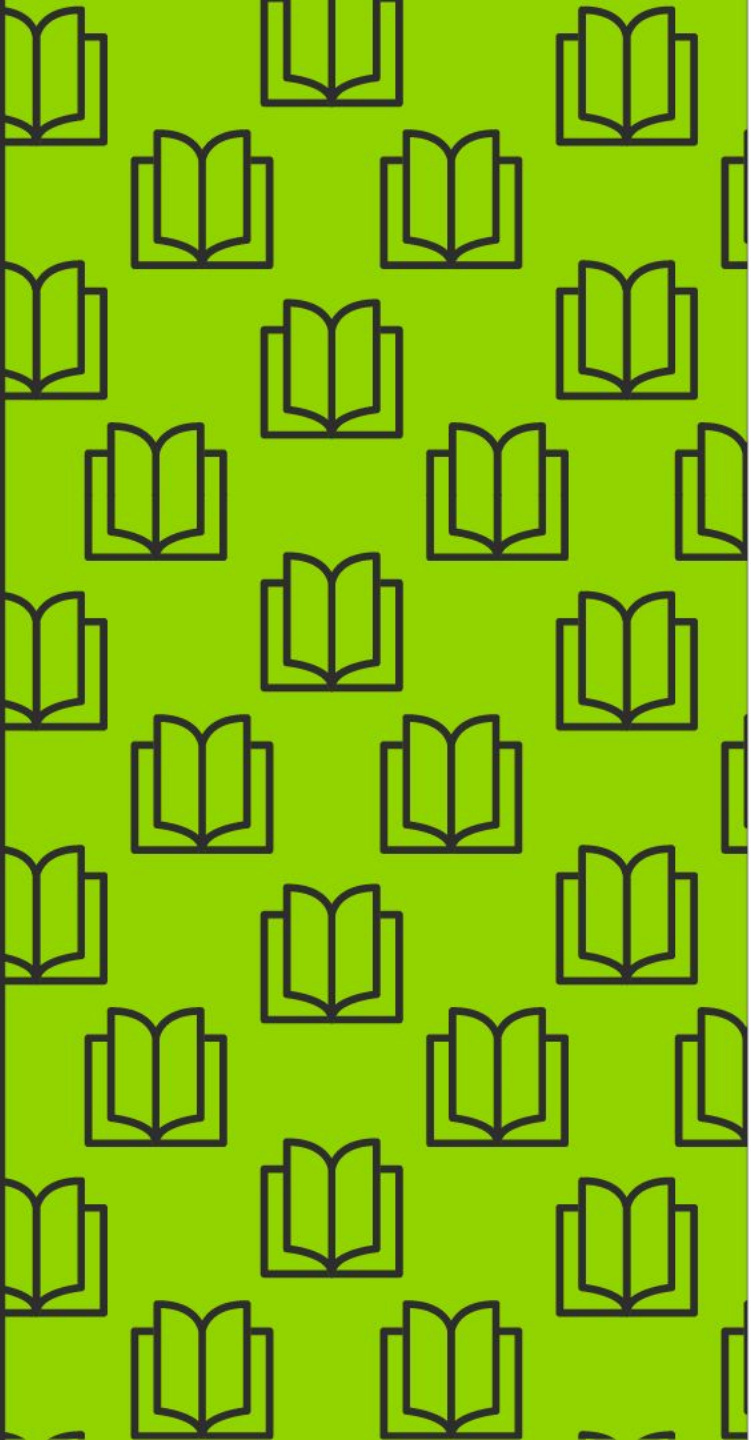
- 1 Admin joins two or more clusters
 - Submariner provides full IP reachability between pods and services among the participating clusters, aka ClusterSet
- 2 Application developers then **export** selected **services** to expose them across the ClusterSet
 - Submariner automatically sets up DNS for the exported services



Step 1 is done once to create cluster-level connectivity. Step 2 (service export) can then be performed on-demand, leveraging the underlying connectivity.


Deployment and management





Research: Optimized Service Discovery Across Clusters

Multi-Cluster Deployment




The diagram shows a globe with several overlapping cloud-like shapes representing different clusters. On the left, a cloud contains a red circle labeled 'Service 1'. On the right, a cloud contains a blue circle labeled 'Service 2'. A large, dark gray rounded rectangle is centered over the globe, containing the text 'Multi-cluster deployment introduces traffic pricing and performance challenges.'.

Multi-cluster deployment
introduces traffic pricing and
performance challenges.

Request Distribution – Round Robin



Solution – Weighted Round Robin



Requirements:

- How to determine the weights?
- Seamless implementation into Kubernetes
- One-hop optimization

Request Distribution – Round Robin

$c \in C$ - Clusters
 $s \in S$ - Services

$cost(c, s)$
 $F(Price, Latency)$


$n_{s,t} \in \mathbb{N}$ (RPS)

Minimize requests cost

$$\text{Min } \sum_{c \in C} \sum_{s \in S} n_{c,s} \cdot cost(c, s)$$

Subject to:

- Requests cannot be



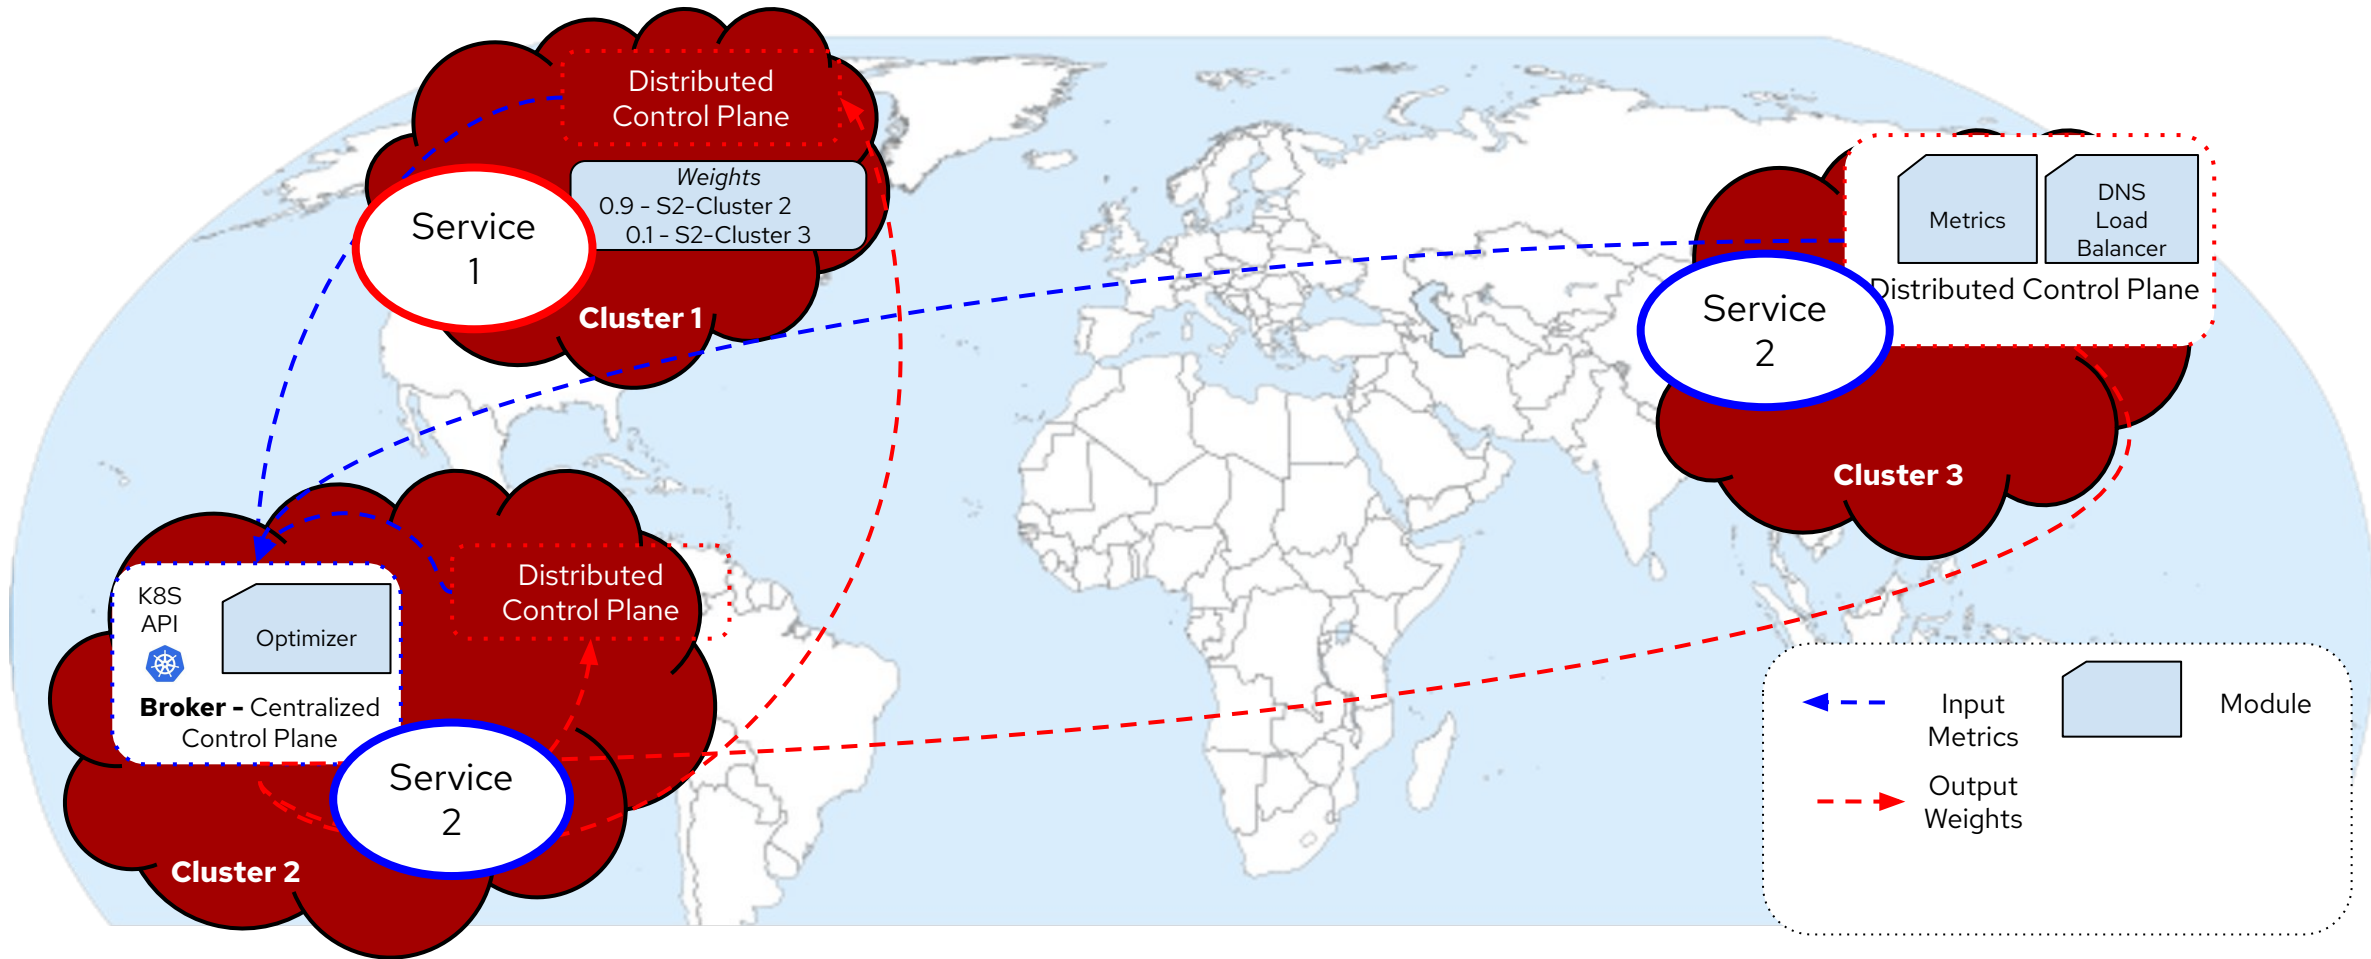
$$= N_c$$

- Target capacity
- Liveness*

$$\sum_{c \in C} n_{c,s} \leq cap(s)$$


$$\sum_{c \in C} n_{c,s} \geq \alpha(c, s)$$

Architecture



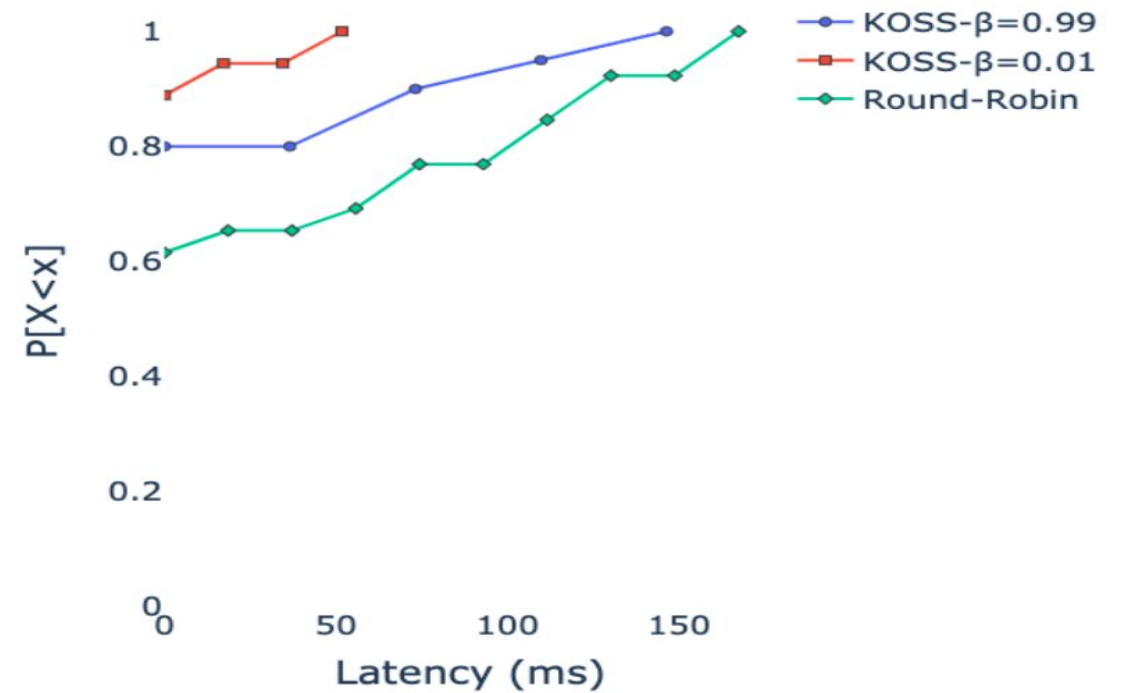
Results

KOSS vs RR pricing min/max β values



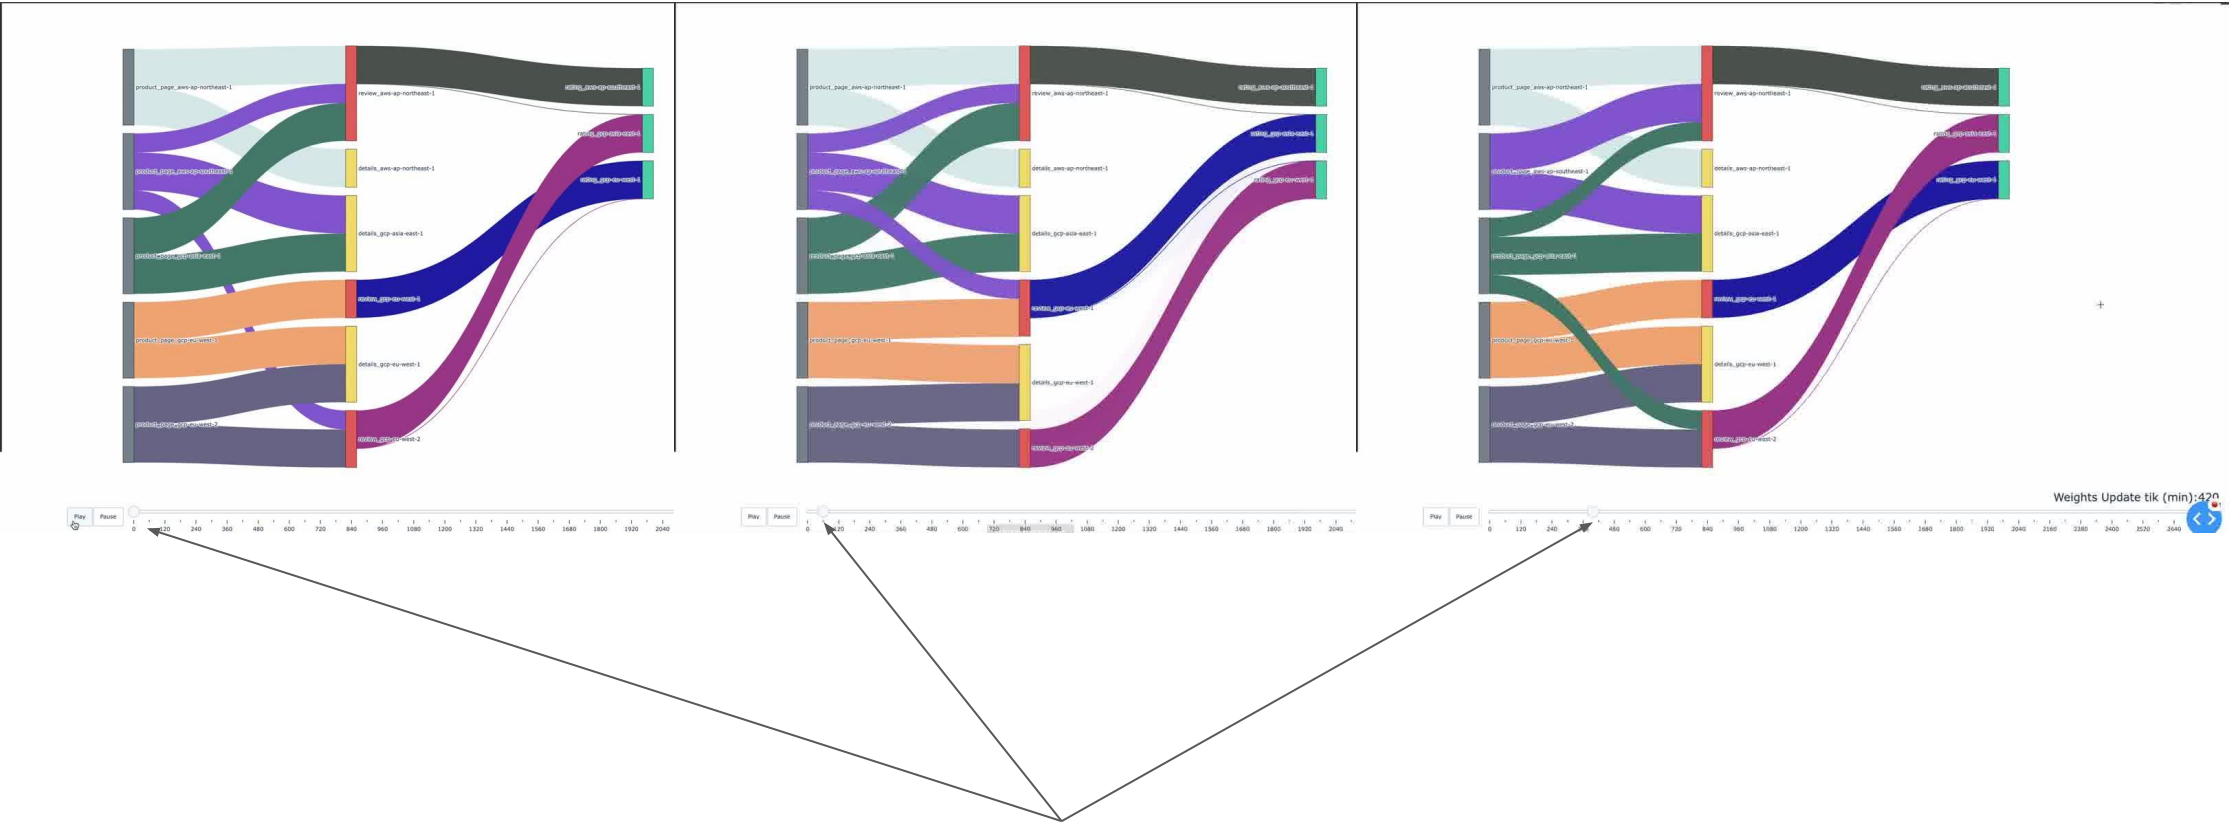
(c)

KOSS vs RR latency max/min β values

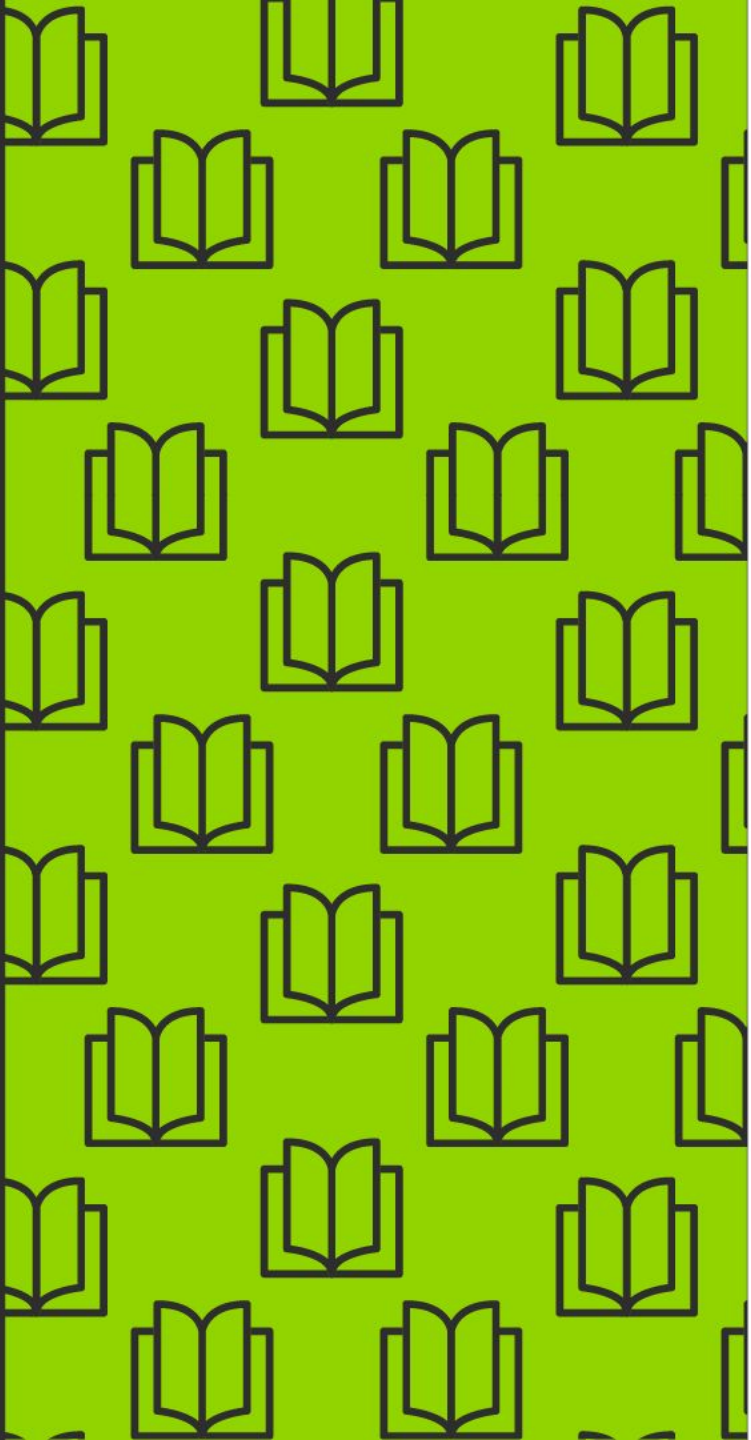


(d)

Results



Adaptive over time and load



Product: Advanced Cluster Management for Kubernetes

Product: Advanced Cluster Management for Kubernetes



Multicluster lifecycle management



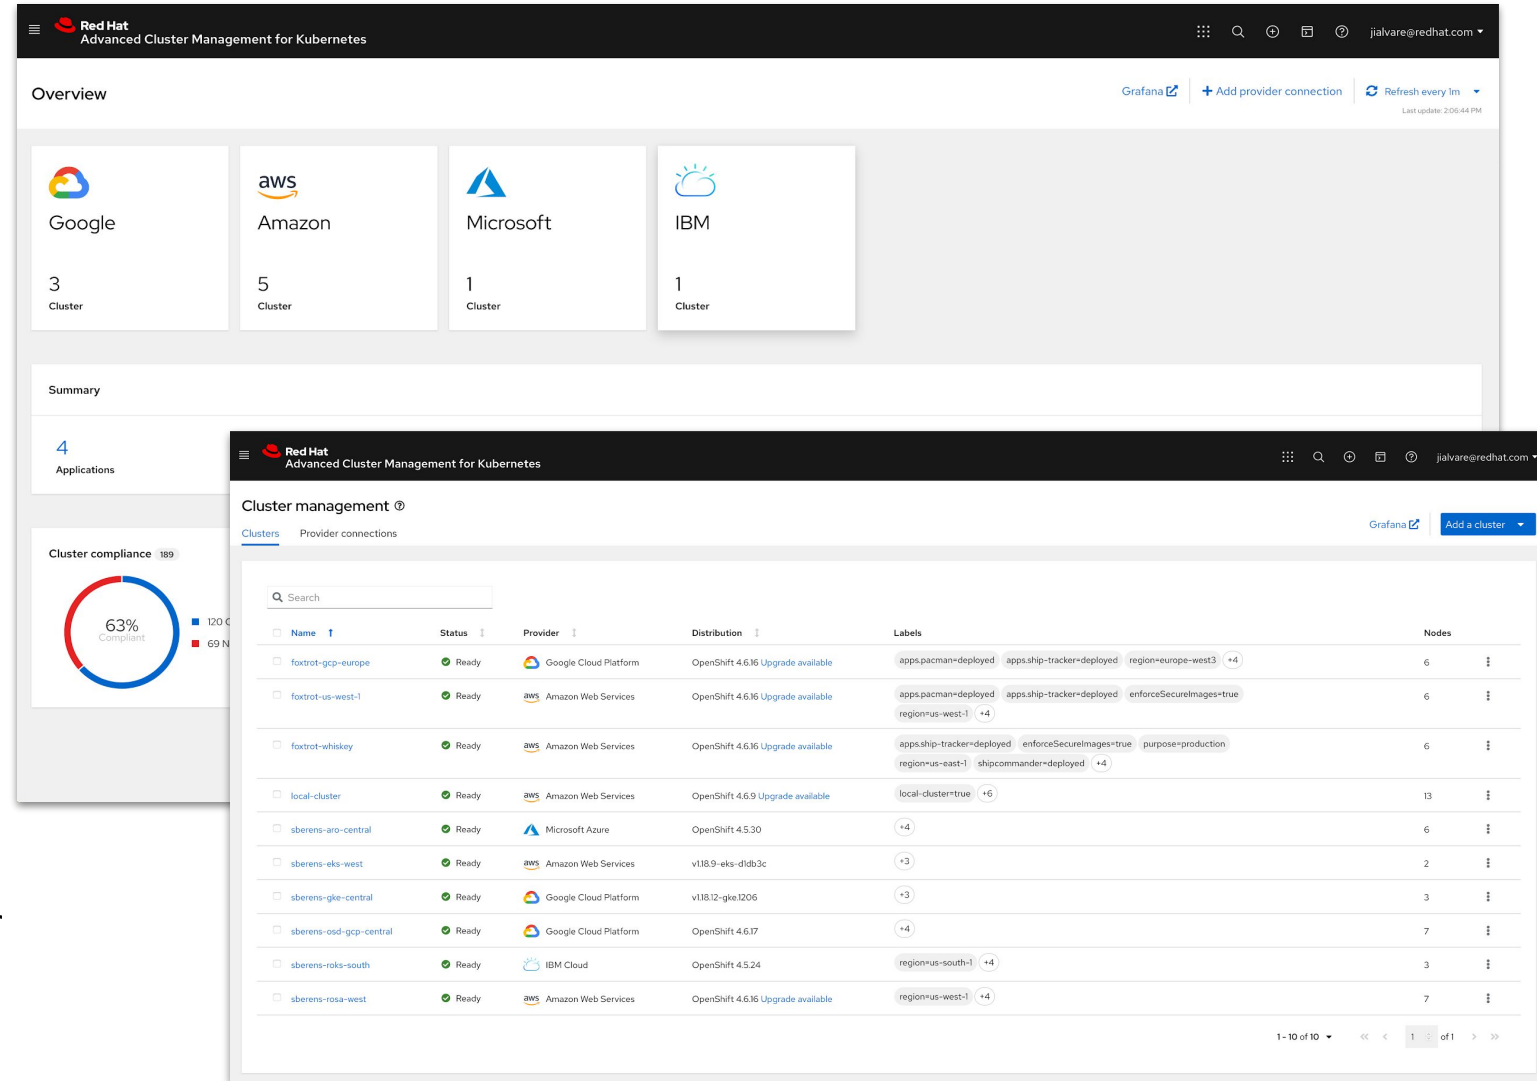
Policy driven governance, risk, and compliance



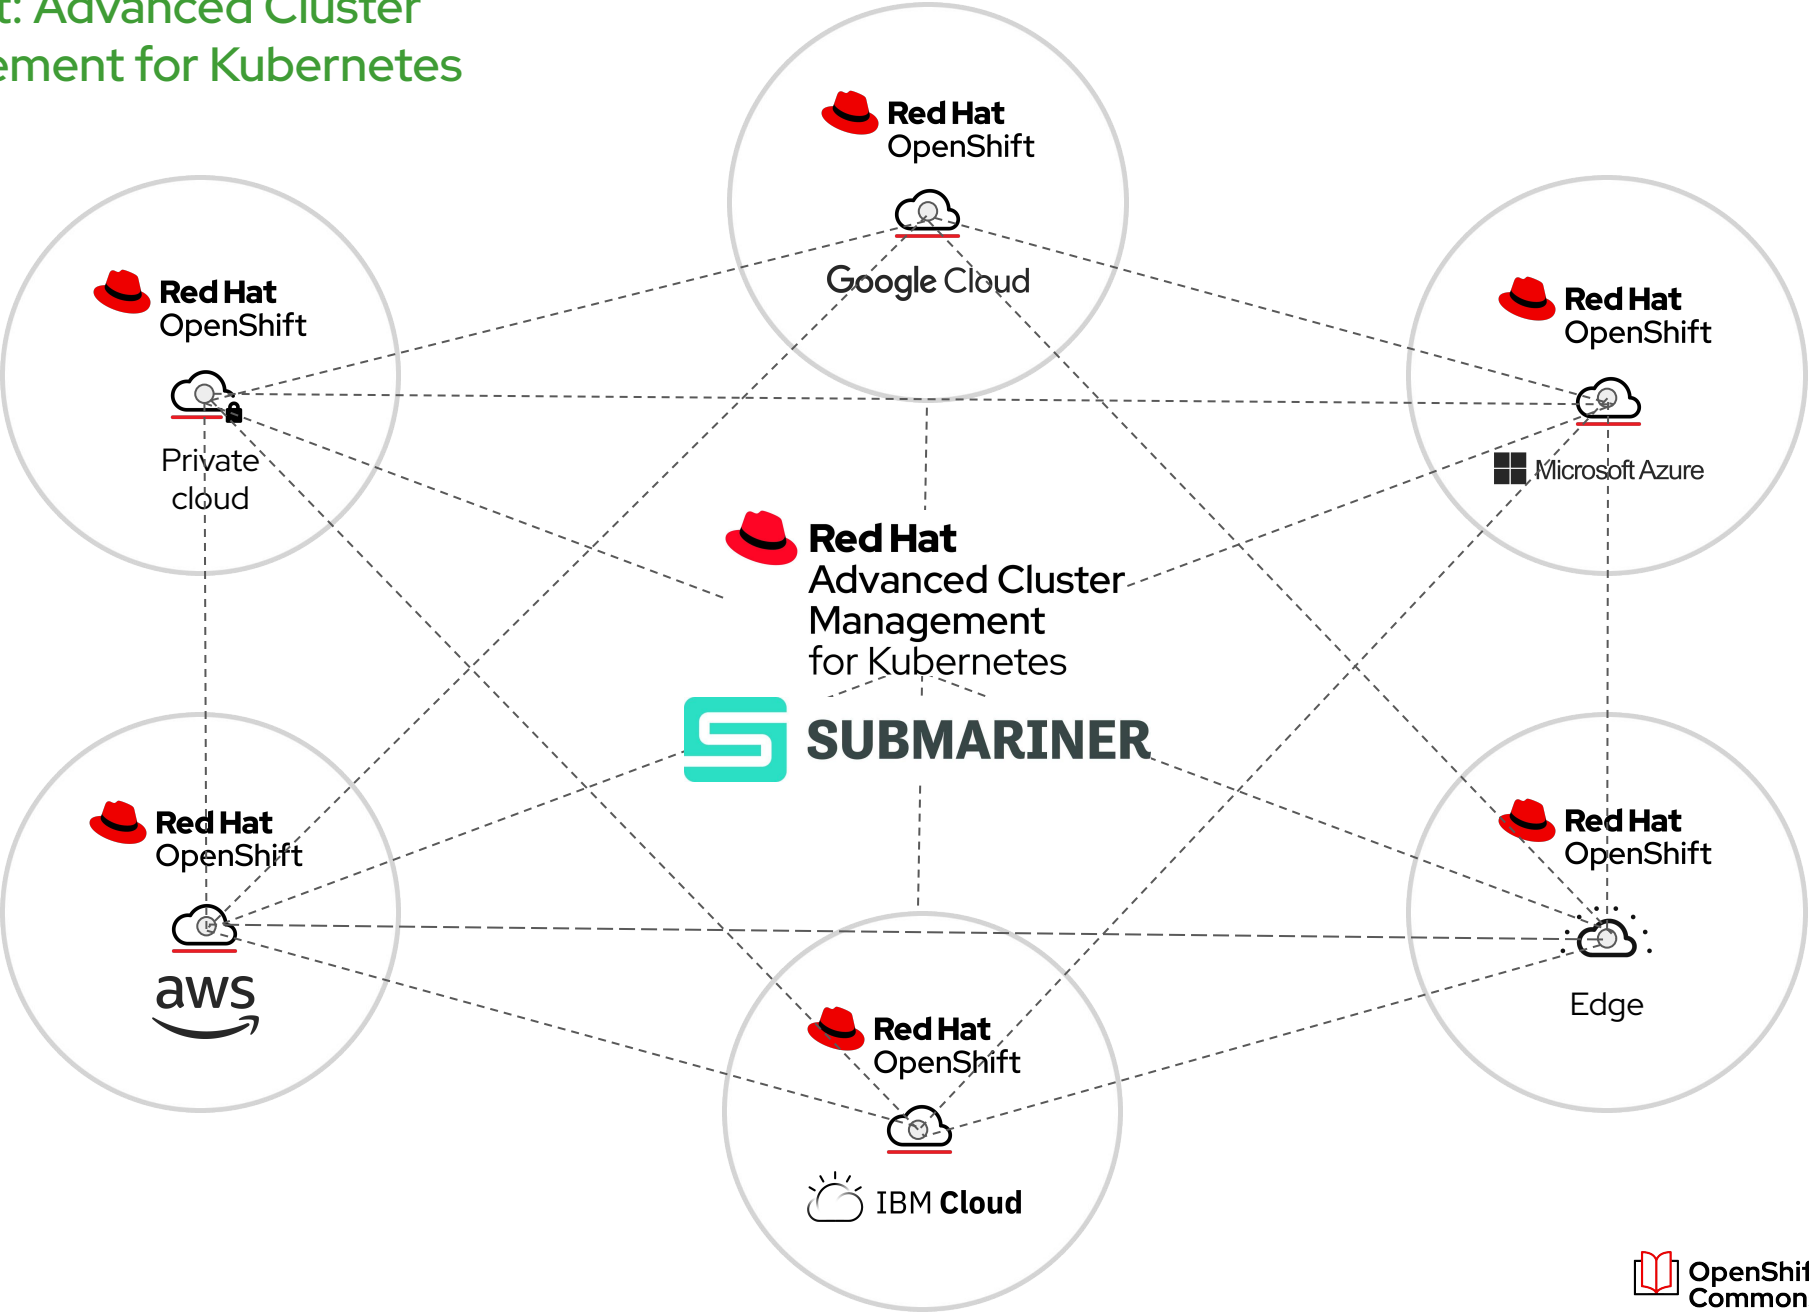
Advanced application lifecycle management



Multicluster observability for health and optimization



Product: Advanced Cluster Management for Kubernetes



Product: Advanced Cluster Management for Kubernetes

Red Hat
Advanced Cluster Management for Kubernetes

Advanced Cluster Management

Home

- Welcome
- Overview

Infrastructure

- Clusters
- Bare metal assets
- Automation
- Infrastructure environments

Applications

Governance

Credentials

End-to-end visibility
[Go to Overview](#)
View system alerts, critical application metrics, and overall system health. Search, identify, and resolve issues that are impacting distributed workloads using an operational dashboard designed for Site Reliability Engineers (SREs).

Cluster lifecycle
[Go to Clusters](#)
Create, update, scale, and remove clusters reliably, consistently using an open source programming model that supports and encourages Infrastructure as Code best practices and design principles.

Application lifecycle
[Go to Applications](#)
Define a business application using open standards and deploy the applications using placement policies that are integrated into existing CI/CD pipelines and governance controls.

Governance, Risk, and Compliance
[Go to Governance](#)
Use policies to automatically configure and maintain consistency of security controls required by industry or other corporate standards. Prevent unintentional or malicious configuration drift that might expose unwanted and unnecessary threat vectors.

Multicluster networking
[Go to Cluster sets](#)
Enable direct networking connection between different on-premises or cloud-hosted Kubernetes clusters by grouping them in cluster sets and enabling the Submariner add-on.

Product: Advanced Cluster Management for Kubernetes

☰

Red Hat

Advanced Cluster Management for Kubernetes

✧ Advanced Cluster Management

Home

Infrastructure

Clusters

Bare metal assets

Automation

Infrastructure environments

Applications

Governance

Credentials

Cluster sets > submariner

submariner

OverviewSubmariner add-onsManaged clustersCluster poolsAccess management

Details

NamesubmarinerNamespace bindings ⓘsubmariner-operator

Multi-cluster network statusⓘ Degraded


Status


3Submariner add-onsGo to Submariner add-ons

3Managed clustersGo to Managed clusters

0Cluster poolsGo to Cluster pools

28

OpenShift Commons

Red Hat OpenShift

9




Resources

Give it a try

- Community project
 - Website: <https://submariner.io>
 - <https://submariner.io/getting-started/quickstart/>
 - <https://submariner.io/operations/usage/>
 - GitHub: <https://github.com/submariner-io>
 - YouTube: <https://tinyurl.com/submariner-youtube>
 - Slack (Kubernetes space): [#submariner](#)
- Red Hat product
 - [ACM+Submariner product documentation](#)

Further reading

- <https://submariner.io/other-resources/>
- Red Hat blogs:
 - <https://cloud.redhat.com/blog/geographically-distributed-stateful-workloads-part-one-cluster-preparation>
 - <https://cloud.redhat.com/blog/geographically-distributed-stateful-workloads-part-two-cockroachdb>
 - <https://cloud.redhat.com/blog/geographically-distributed-stateful-workloads-part-3-keycloak>
 - <https://cloud.redhat.com/blog/geographically-distributed-stateful-workloads-part-four-kafka>
 - <https://cloud.redhat.com/blog/geographically-distributed-stateful-workloads-part-five-yugabytedb>



Thank you

Commons builds connections and collaboration across OpenShift communities, projects and stakeholders. In doing so we'll enable the success of customers, users, partners, and contributors as we deepen our knowledge and experiences together.

 commons.openshift.org

 facebook.com/openshift

 youtube.com/OpenShift

 twitter.com/openshift