


---

# Submariner: Multi Cluster Networking for Kubernetes


Nir Yechiel, Red Hat  
@nyechiel


Developer Community TLV  
December 17, 2020






Kubernetes Admin





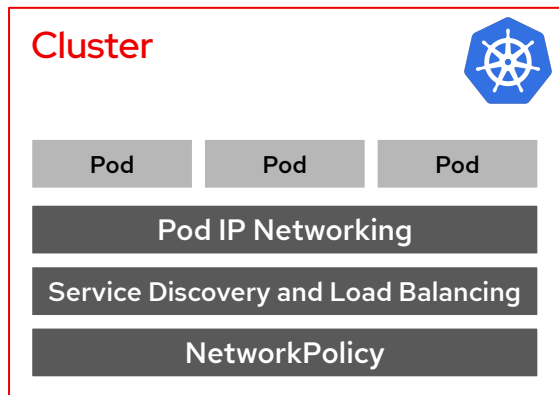





# What is Submariner?

- Enables direct network connectivity between multiple Kubernetes clusters
- Exposes a set of new custom resources (CRDs) backed by the Kubernetes datastore
- Open source, vendor neutral project: <https://submariner.io/>
- Can be deployed using Operator, Helm and subctl

# Kubernetes cluster networking



# Multi cluster networking with Submariner



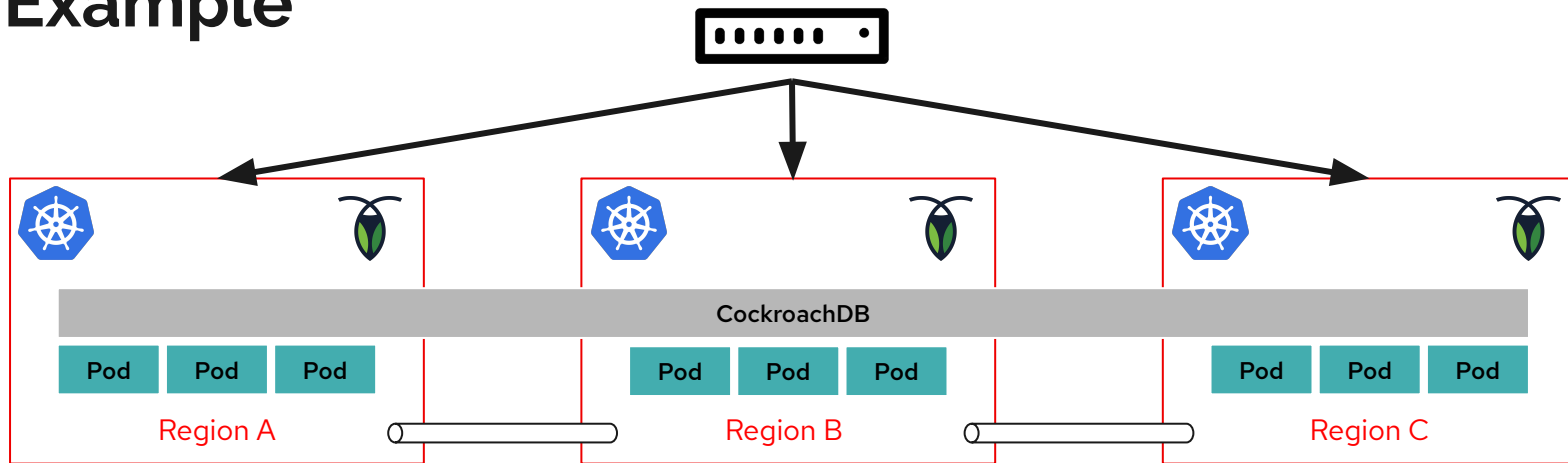


## Common use cases

- Application availability
- Disaster Recovery (DR)
- Data residency guidelines



# Example



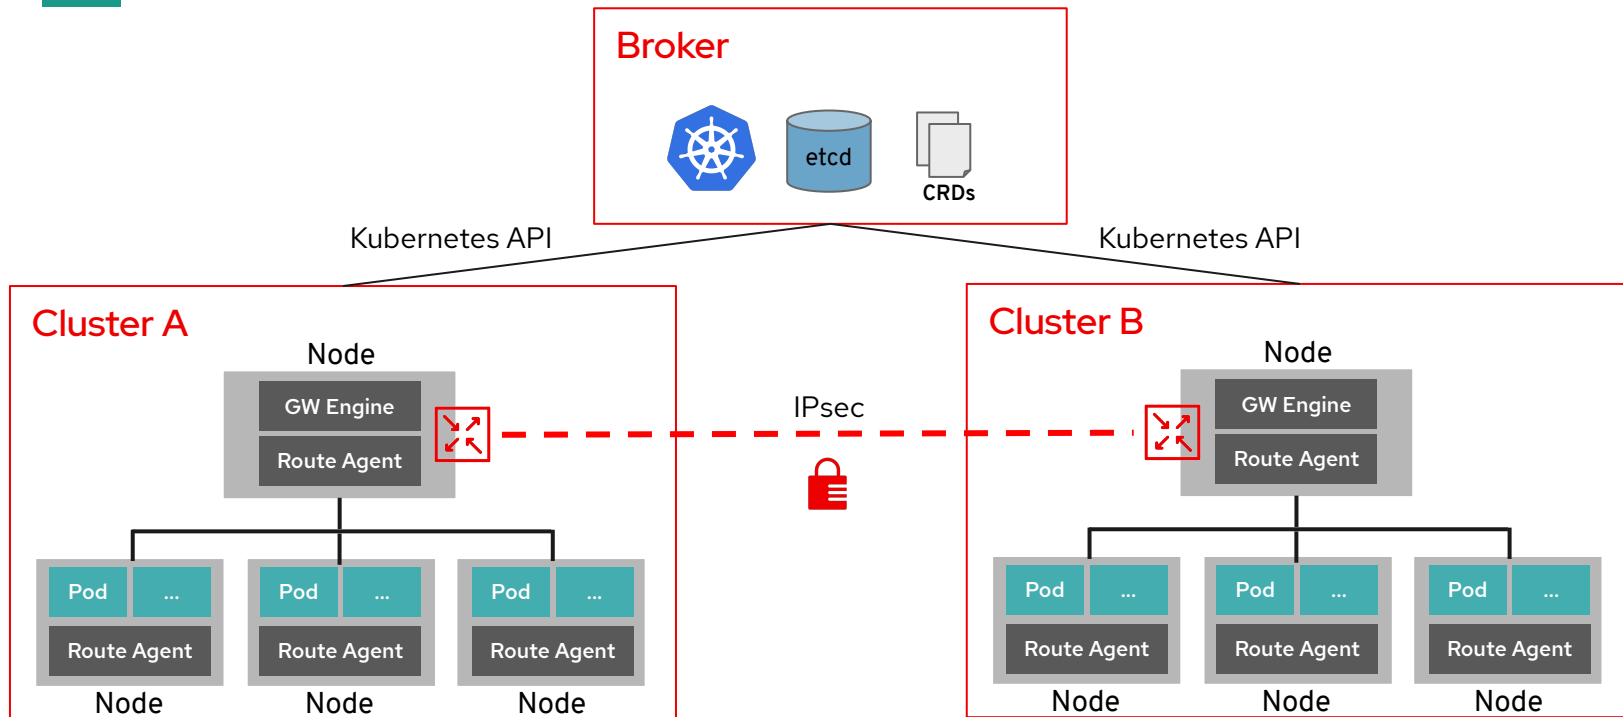
- Example: multi-region CockroachDB cluster
- Kubernetes clusters in multiple regions, with replicas of the same service running in each cluster
- Goal: keep data close to the user, while reducing latency and fault tolerance to improve user experience



# Benefits

- Direct east/west Pod-to-Pod and Pod-to-Service routing across clusters
  - ... without requiring proxies or exposing your backend services to the internet
- Works with any Kubernetes cluster, compatible with different network (CNI) providers
- Deploy services across clusters
  - Beyond connectivity, also address Service Discovery (Lighthouse) and NetworkPolicy (Coastguard)
- Enhanced security
  - Cross-cluster traffic is encrypted with IPsec by default; WireGuard also supported

# High level architecture



# Service Discovery (DNS)

- Lighthouse is an implementation of the Kubernetes [Multi-Cluster Service API](#)
- **clusterset** - a group of clusters with a high degree of mutual trust that share services
  - Within a `clusterset`, all namespaces with a given name are considered to be the same
- **ServiceExport** (CRD) - used to specify which services should be exposed across clusters
  - To export a service, a ServiceExport should be *explicitly* created within the cluster
  - When a ServiceExport is created, this will cause a domain name of `<service>.<ns>.svc.clusterset.local` for the multi-cluster service to become accessible
- **ServiceImport** (CRD) - in-cluster representation of a multi-cluster service

# subctl

- A command-line utility to ease deployment and verification of Submariner

```
# subctl join --kubeconfig output/kubeconfigs/config-cluster1 broker-info.subm
```

```
# subctl join --kubeconfig output/kubeconfigs/config-cluster1 broker-info.subm
```

---

# Try Submariner

- Website: <https://submariner.io>
  - [https://submariner.io/getting\\_started/quickstart/](https://submariner.io/getting_started/quickstart/)
- GitHub: <https://github.com/submariner-io>
- YouTube: <https://tinyurl.com/submariner-youtube>
- Slack (Kubernetes space): [#submariner](#)